

TEMPLATE FOR ASSIGNMENT #4 REPORT

Student Name and CCID:

- A. Sidharth Srinivasan_____ (Name) sr11_____ (CCID)
B. James Schaefer-Pham_____ (Name) jschaefer_____ (CCID)
C. Adit Rada_____ (Name) rada_____ (CCID)
-

By submitting this assignment the students named above confirm that they have worked on it themselves without any help by other people. If any external resources were used please state which ones and how they were used:

PART 1

Task A (1) (no index):

| Cardinality of Table Parts | Average Processing time for index free Q1 (ms) |
|----------------------------|--|
| 100 | 0.262ms (0.0002628850936889648) |
| 1000 | 0.360ms (0.00035992860794067385) |
| 10,000 | 0.356ms (0.00033558368682861327) |
| 100,000 | 0.429ms (0.0004287815093994141) |
| 1,000,000 | 0.556ms (0.0005563831329345703) |

| Cardinality of Table Parts | Average Processing time for index free Q2 (ms) |
|----------------------------|--|
| 100 | 0.368ms (0.00036847829818725586) |
| 1000 | 0.370ms (0.0003700733184814453) |
| 10,000 | 4.45ms (0.004479818344116211) |
| 100,000 | 39.1ms (0.039133446216583254) |
| 1,000,000 | 379.1ms (0.3791451859474182) |

Task B (2):

Compare, contrast and explain the trends observable in both tables above (Task A)

The average time increases as the size of the database increases for both queries, this is consistent with the fact that as the no. of records increases, the query needs to iterate over more tuples and hence longer time.

The average time for executing query 2 is higher than that for query 1 for the corresponding instance of the database. This is as since needsPart is not the unique (unlike partNumber which was the primary key), the query must process over the entire database to find all the records with that needsPart no.; however, query 1 could stop iterating over the database once the record is found hence it would scan, on average, about $n/2$ records; while query 2 scans all n records.

Task C (3) (using index):

| Cardinality of Table Parts | Average Processing time for indexed Q1 (ms) |
|----------------------------|---|
| 100 | 0.110ms (0.00011001348495483399) |
| 1000 | 0.112ms (0.00011253118515014649) |
| 10,000 | 0.191ms (0.00019063711166381837) |
| 100,000 | 0.190ms (0.00018982887268066407) |
| 1,000,000 | 0.291ms (0.00029119968414306643) |

| Cardinality of Table Parts | Average Processing time for indexed Q2 (ms) |
|----------------------------|---|
| 100 | 0.0898ms (0.00008979082107543946) |
| 1000 | 0.134ms (0.0001339411735534668) |
| 10,000 | 0.172ms (0.00017249822616577148) |
| 100,000 | 0.180ms (0.00017951250076293944) |
| 1,000,000 | 0.242ms (0.00024185895919799803) |

Task D (4):**Compare, contrast and explain the trends observable in both tables above (Task C)**

The average time increases as the size of the database increases for both queries. The times for both queries on corresponding instances of the databases is about the same despite the fact that query 2 should actually take longer as it scans all tuples in the relation whereas query 1 scans about $n/2$ tuples. The fact that both have similar times suggests that

the index helps query 2 more than it helps query 1. This makes sense as the index is on needsPart and we are doing an equality search on needsPart in query 2; however, we are doing equality search on partNumber in query 1, so index unlikely to help.

Task E (5):

Compare, contrast and explain the trends observed in Task D to the trends observed in Task B. Discuss the cost-benefit of the index space cost and query performance.

Space before and after:

A4v100: 8kb -> 12kb
A4v1k: 36kb -> 60kb
A4v10k: 248kb -> 428kb
A4v100k: 2464kb -> 4224kb
A4v1M: 25,484kb -> 43,056kb

The average times are less with the index for both queries, though the times are only slightly less for query 1 but are significantly less for query 2.

For e.g., for query 2 on Av1M the average time reduces from 379.1ms to 0.242ms(Almost 100% reduction in time). Though we must note that the space needed, on average, for all the database instances increases by about 1.7 times.

We can conclude, that the benefits of an index are not worth the space costs for query 1 but are more than worth it (due to the enormous time savings) for query 2.

PART 2

Task F (6) (no index):

| Cardinality of Table Parts | Average Processing time for index-free Q3 (ms) |
|----------------------------|--|
| 100 | 0.310ms (0.0003096318244934082) |
| 1000 | 1.510ms (0.0015097546577453613) |
| 10,000 | 9.825ms (0.009825007915496826) |
| 100,000 | 115.1ms (0.11507606744766236) |
| 1,000,000 | 1611.7ms (1.6116997051239013) |

Task G (7) (using index):

| Cardinality of Table Parts | Average Processing time for indexed Q3 (ms) |
|----------------------------|---|
| 100 | 0.289ms (0.00028920888900756835) |
| 1000 | 1.418ms (0.0014175820350646974) |
| 10,000 | 9.462ms (0.009461996555328369) |
| 100,000 | 250.0ms (0.25021379232406615) |
| 1,000,000 | 7643.7ms (7.6436997051239013) |

Task H (8):

Compare, contrast and explain the trends observed in Task F to the trends observed in Task G. Discuss the cost-benefit of the index space cost and query performance.

For both tests, as the size of the database increases, the average time increases.

It seems like the index does not make any difference in terms of time. For e.g., for A4v100k the time increased from 115ms to 250ms with the index.

So, the index is actually making the performance worse (at least on my machine).

Thus, we can say that the index is probably not worth the increased space costs. One hypothesis for the time not changing is that the time taken to iterate over all rows when calculating the average dominates the total time and using an index does not speed up this process anyway. So, looking up in the index and then iterating over all the rows is more time consuming than iterating over all the rows.

PART 3**Task I (9) (no index):**

| Cardinality of Table Parts | Average Processing time for no-index Q4 (ms) |
|----------------------------|--|
| 100 | 0.07976531982421875 |
| 1000 | 0.1404428482055664 |
| 10,000 | 0.8568572998046875 |
| 100,000 | 8.481192588806152 |
| 1,000,000 | 77.7010440826416 |

Task J (10):

Define an index that you believe will optimize Q4 and explain why you think so.

```
SELECT p.partNumber
FROM Parts p
WHERE p.madeIn = :C
AND p.partPrice = (
SELECT MAX(q.partPrice)
FROM Parts q
WHERE q.madeIn = :C)
```

:C is the country code

We used the command **CREATE INDEX MaxCost On Parts (madeIn, partPrice)** to make or index. In the aggregate query, we only need to concern ourselves with parts that have a given country code, so having quick access to parts with that value of **madeIn** speeds the query up. The same is true of the greater query. We also include **partPrice** since the aggregate query gives us a set value for **partPrice** that we can lookup.

Task K (11) (using index):

| Cardinality of Table Parts | Average Processing time for indexed Q4 (ms) |
|----------------------------|---|
| 100 | 0.03988742828369141 |
| 1000 | 0.040040016174316406 |
| 10,000 | 0.06110668182373047 |
| 100,000 | 0.10442733764648438 |
| 1,000,000 | 0.19931793212890625 |

Task L (12):

Compare, contrast and explain the trends observed in Task K to the trends observed in Task I. Discuss the cost-benefit of the index space cost and query performance.

Size without an index:

A4v100: 8 KB
A4v1k: 36 KB
A4v10k: 248 KB
A4v100k: 2 464 KB
A4v1M: 25 484 KB

Size with an index:

A4v100: 12 KB
A4v1k: 56 KB
A4v10k: 408 KB
A4v100k: 4 028 KB
A4v1M: 41 120 KB

It appears that databases with an index use about 1.55 times as much memory as a database without indices from our samples on average. It seems that the memory used grows at a slightly faster than linear rate. Growth in memory cost seems to be negligible however, as the cost add an index to a database with 100 entries is an increase of 1.5 times, as opposed to 1.6 times for a database with 1000 times more entries (1 million entries).

In terms of performance, including an index provides very obvious improvements. The smallest improvement, as seen in the 100 entry database, doubles the query speed. The speedup caused by the index becomes even larger for larger databases. By the time the database size reaches 1 million, adding an index makes the query speed 390 times faster. From this, it seems that adding an index to a small database may not necessarily be worth the increase in memory, though this depends primarily on if the user prioritizes performance or memory. However, for large databases, it is almost certainly worth it to use the index since it greatly decreases the query time with relatively little impact to the database size.

PART 4

Task M (13) (no index):

| Cardinality of Table Parts | Average Processing time for index-free Q5 (ms) |
|----------------------------|--|
| 100 | 1.2758779525756836 |
| 1000 | 73.99272918701172 |
| 10,000 | 7 752.774424552918 |
| 100,000 | N.A. |
| 1,000,000 | N.A. |

Task N (14) (no index):

| Cardinality of Table Parts | Average Processing time for index-free Q6 (ms) |
|----------------------------|--|
| 100 | 0.2506828308105469 |
| 1000 | 0.8909940719604492 |

| | |
|-----------|-------------------|
| 10,000 | 11.44637107849121 |
| 100,000 | 169.4697141647339 |
| 1,000,000 | 5 344.4664478302 |

Task O (15):

Compare, contrast and explain the trends observed in Task M to the trends observed in Task N

We can clearly see from the details from the two tasks that Query 6 is much faster than Query 5. On trying to run Query 5 on 100,000 and 1,000,000 Cardinality, we will notice that the time runs on $O(n^2)$. Every single run takes about 10 minutes for 100k in case of query 5 and hence had to be avoided.

On the other hand, Query 6 is definitely much faster in implementation, with an average of just approximately 5.3 seconds for 1M cardinality, in contrast to the 7.75s for just 10k cardinality

Task P (16):

Define an index that you believe will optimize Q6 and explain why you think so

```
SELECT COUNT(1)
FROM Parts P
WHERE P.partNumber not in (select Q.needsPart from Parts Q);
```

We used the command **CREATE INDEX idxneedParts On Parts (needParts)** to make our index. We used this because it allows us to quickly lookup **partNumber** in the **needsPart** indexed to see if there is a part that has it as a **needsPart** value.

Task Q (17) (with index):

| Cardinality of Table Parts | Average Processing time for indexed Q6 (ms) |
|----------------------------|---|
| 100 | 0.008619785308837891 |
| 1000 | 0.4877700805664062 |
| 10,000 | 6.729097366333008 |
| 100,000 | 56.74569606781006 |
| 1,000,000 | 2 817.062759399414 |

Task R (18):

Compare, contrast and explain the trends observed in Task N to the trends observed in Task Q. Discuss the cost-benefit of the index space cost and query performance.

Size without an index:

| | |
|----------|-----------|
| A4v100: | 8 KB |
| A4v1k: | 36 KB |
| A4v10k: | 248 KB |
| A4v100k: | 2 464 KB |
| A4v1M: | 25 484 KB |

Size with an index:

| | |
|----------|-----------|
| A4v100: | 12 KB |
| A4v1k: | 60 KB |
| A4v10k: | 428 KB |
| A4v100k: | 4 224 KB |
| A4v1M: | 43 056 KB |

Similar to the previous indices, including the index caused the database to grow by a factor of less than 2 in all cases. In this instance, the most memory growth is in the A4v1M database which grew by a factor of 1.69.

We are able to notice that the use of an index has drastically reduced performance time. The time taken for operations on 1M Cardinality now took 2.5s lesser which is a major improvement. On average, query time is cut in half when using an index.

We can easily say that the cost-benefit is much higher than the space cost of the index. Unless we prioritize memory much more than performance, it is better to use an index. However, for very large databases (cardinality 1 billion+) it may be better to not use an index, since it provides only linear performance improvement whereas memory growth is faster than linear. This can be attributed to the ease in iterating over the indices as compared to everything.